

A judge's calibration does not carry, least of all across training that targets the label

A rubric judge's recall and false-alarm rate are measured on one population and used to bound the refusal rate of another. Left: how often that bound falls below the truth, one dot per wording of the judge's rubric. Right: what happened to the judge's recall.

- recall where the calibration was measured
- where it was applied (both panels)

