A judge's calibration does not carry across prompt populations, nor, in three seeds of one run, across training that targets the label
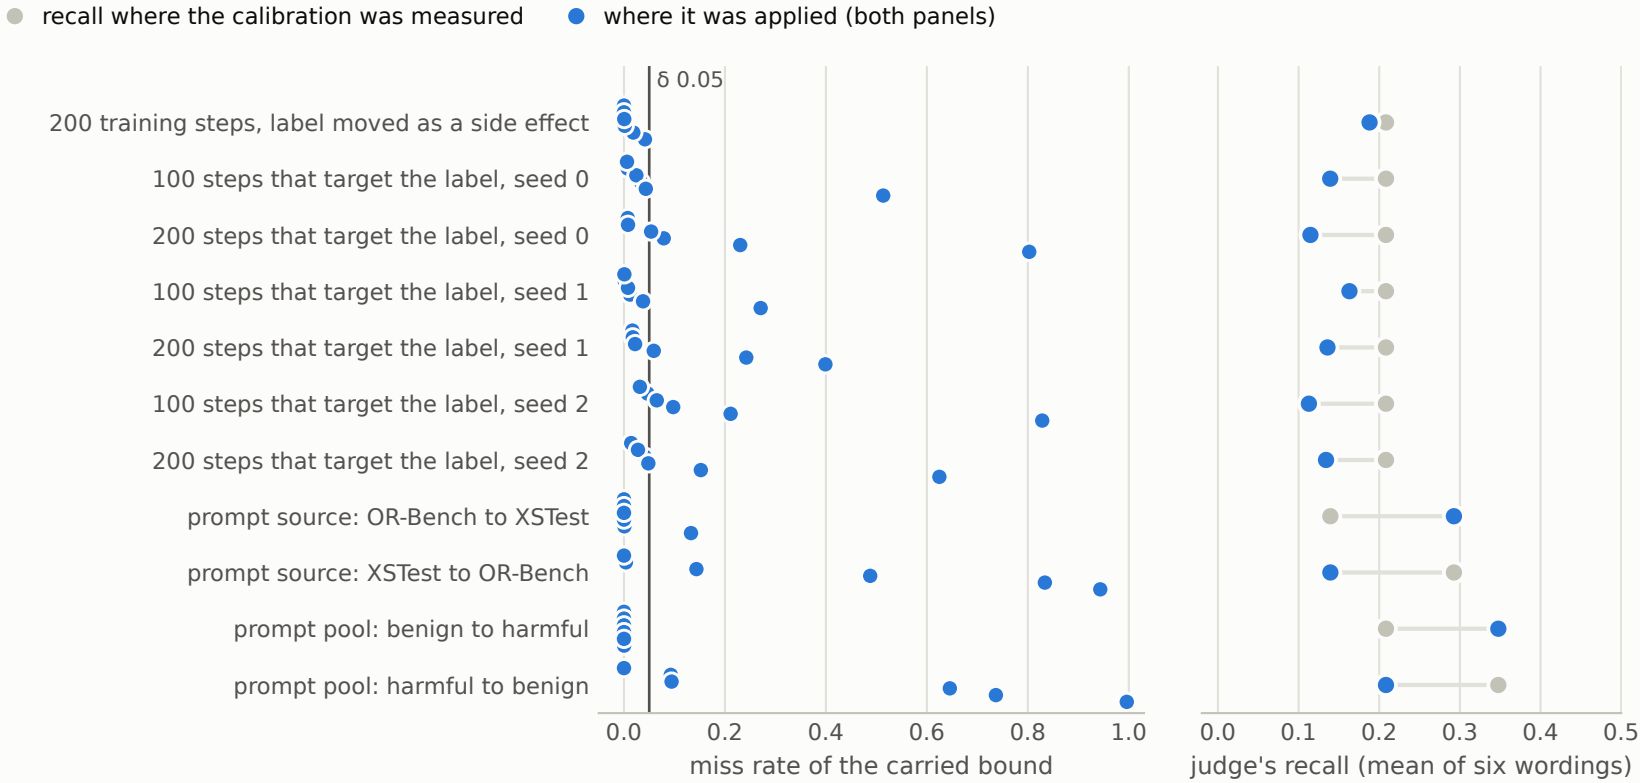
A rubric judge's recall and false-alarm rate are measured on one population and used to bound the refusal rate of another. Left: how often that bound falls below the truth, one dot per wording of the judge's rubric. Right: what happened to the judge's recall.