A rubric judge's calibration, carried to another prompt population, fails in at least one direction of each pair, and fails across training that targets the label in three seeds of one run
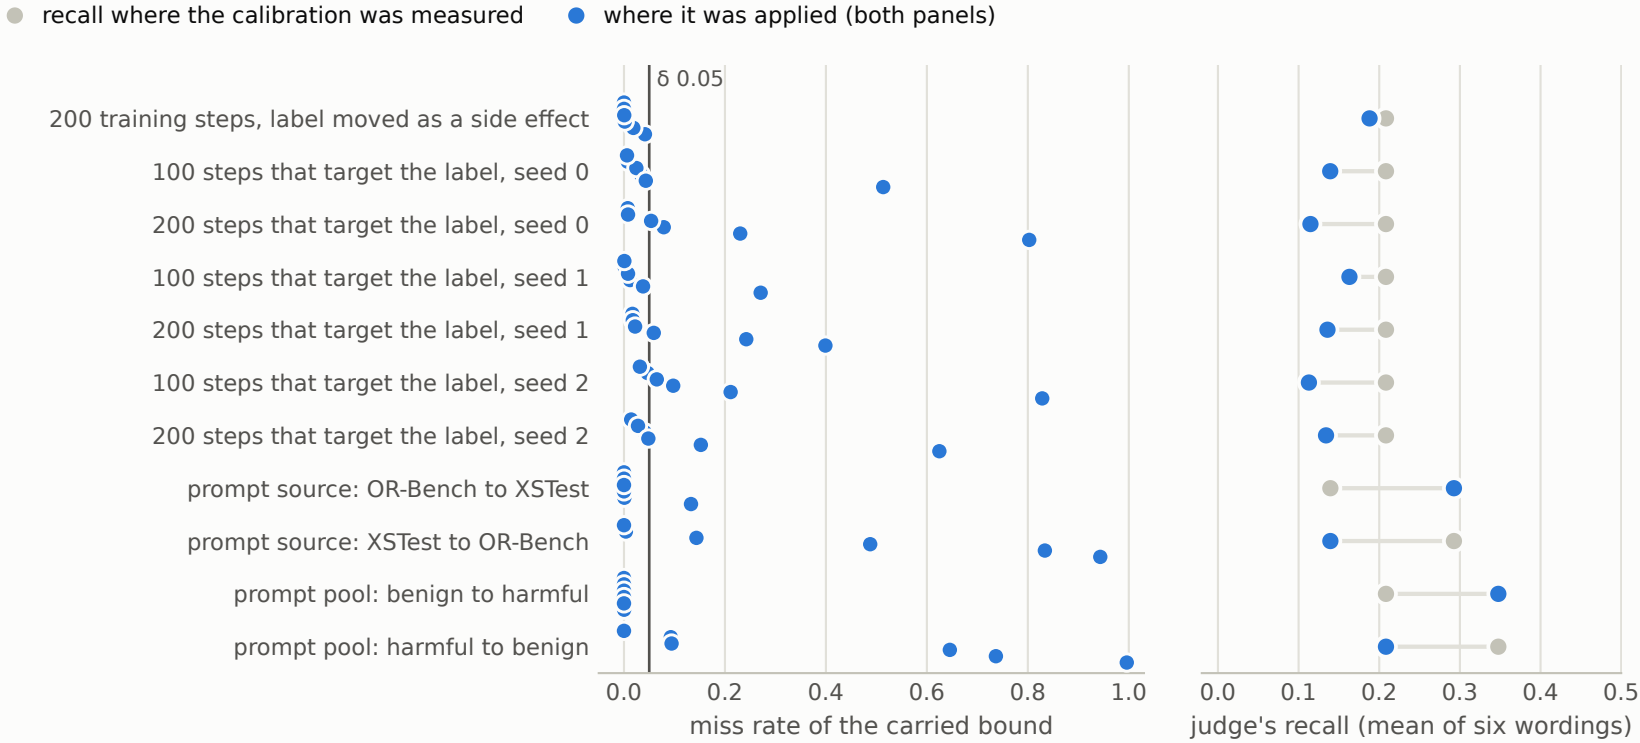
A rubric judge's recall and false-alarm rate are measured on one population and used to bound the refusal rate of another. Left: how often that bound falls below the truth, one dot per wording of the judge's rubric. Right: what happened to the judge's recall.