

In human terms the two policies refuse equally often, and less often than the guard says

220 responses labelled by one annotator, a stratified sample read with its design weights; bars are conservative 95% intervals. The guard's rate is its share of the whole pool, so it has no interval.

● reference model ● policy trained under the constraint

