

Certifying behaviour rates of language-model policies: what holds, what a label buys, and what does not carry

Draft of 2026-10-06, not formatted for any venue

Abstract

A Seldonian algorithm (Thomas et al., 2019) returns a solution only when a high-confidence test on held-out data supports a stated constraint, and otherwise returns “no solution found”. Used to post-train a language model, its last step, the safety test, is a one-sided confidence bound on the rate of a judged behaviour. We take that test on its own, as a certificate for a fixed policy, and audit it against known truths: synthetic environments, resampled real responses, and published traces of frontier models. No exact bound was over its level in any cell. The normal-quantile intervals of PPI++ and StratPPI miss in up to 24% of draws at a nominal 5%; a bootstrap-t limit is over its level in no cell when labels are allocated in proportion to stratum size. Stratifying the safety set by the reference model’s own per-prompt rate multiplies the effective sample by 1.4 to 5.3 for labels at rates of 9% to 66%, under a strata rule chosen on the same data. Three things do not carry: a judge’s calibration, across prompt populations or across training that targets the label; independent-sample bounds on crossed benchmark designs; and a stratified labelling sheet read as a random sample. On AgentDojo’s published runs the usual per-pair bound is over its level for 27 of 28 pipelines, and one of 28 certifies a 5% attack success rate under a bound clustered by user task. A certificate of a trained policy in human labels is costed and prepared, not run. A normal-approximation budget at the measured difference of about zero puts a 2-point refusal margin at about 420 labelled prompt pairs for an even chance of passing; in a check of the design with synthetic labels, 300 pairs certify a margin of 3.7 points with an approximate limit and 6.5 with the exact one.

1 Introduction

Reports on language-model policies state rates: how often a model refuses a benign request, how often a tool-using agent follows an injected instruction, how often a robot controller trips a safety stop. A decision to deploy needs a different object: a statement that the rate is at most some threshold, together with the chance that the statement is wrong. We call that statement a certificate.

The Seldonian framework (Thomas et al., 2019) produces this object as the last step of a learning algorithm. The algorithm searches for a solution on one part of the data, tests it on a part the search never saw, and returns it only if a one-sided confidence bound at level δ puts the constrained quantity under its threshold. Otherwise it returns “no solution found” (NSF). Sections 2 to 4 set out the algorithm, its use to post-train a language model, and the theory of each step.

Three features of language-model evaluation can each void the guarantee without any visible sign. The label is not observed: a guard model, a rubric judge or a benchmark harness produces it, and its errors differ from one population of responses to the next. The sample is rarely independent: benchmarks cross tasks with attacks, and failures cluster. And the policy has often been trained against the label that is then used to certify it. The rates of interest are also small, which is where normal approximations are known to fail (Bowyer et al., 2025).

This paper does not propose a new way to train. It takes the safety test, applies it to policies that are fixed by the time they are tested, and asks when its guarantee holds for a language model, what it costs in labels, and what breaks it. Section 5 states the proposal, section 6 the set-up of the evaluation, and sections 7 to 10 the results.

2 The Seldonian algorithm

A standard learning algorithm returns the parameters θ that score best on an objective, and promises nothing about a constraint that entered as a penalty. A Seldonian algorithm (Thomas et al., 2019) starts from the promise. The user supplies constraint functions g_i , with $g_i(\theta) \leq 0$ meaning the behaviour is acceptable, and a tolerance δ_i for each. The algorithm a maps data D to a solution or to NSF, and must satisfy

$$\mathbb{P}(g_i(a(D)) \leq 0) \geq 1 - \delta_i \quad \text{for every } i,$$

where $g_i(\text{NSF}) = 0$ by definition, so declining never counts as a violation. With $\delta = \sum_i \delta_i$, the probability of returning a solution that violates some constraint is at most δ . The standard construction has three steps.

1. **split** D into a candidate set D_c and a safety set D_s
2. **candidate** using D_c only, search for the θ_c with the best objective
 selection among those predicted to pass step 3
3. **safety** using D_s only, once: compute U_i , a $(1 - \delta_i)$ upper confidence
 test bound on $g_i(\theta_c)$, from unbiased estimates $\hat{g}_i$ of it
 return θ_c if every $U_i \leq 0$, otherwise return NSF

It works because θ_c is a function of D_c alone. Given D_c it is one fixed solution and D_s is a fresh sample, so step 3 is an ordinary confidence bound and falls below the true $g_i(\theta_c)$ with probability at most δ_i . The search may look at any number of candidates, because exactly one is tested, once.

The guarantee is on the output. It does not depend on the search converging or on the prediction in step 2 being accurate. Nothing is promised about how often a solution is returned: an algorithm that always returns NSF is Seldonian, so the rate at which solutions are returned is reported as a measure of quality. And the guarantee is exact only if the bound is. With a normal-approximation bound such as Student’s t , Thomas et al. call the algorithm quasi-Seldonian.

3 Seldonian post-training of a language model

Post-training starts from a released, instruction-tuned model and changes its behaviour with a small number of further updates. In the work this paper builds on, the updates are reinforcement learning on a reward (GRPO, Group Relative Policy Optimisation) applied to a LoRA adapter, with the base weights frozen. The untrained base model is the *reference policy*. The concern is that optimising a reward model’s score moves behaviours the reward does not measure, such as how often a benign request is refused. Table 1 gives each part of section 2 its language-model form, as implemented in `seldonian/llm/`.

Table 1. The Seldonian algorithm instantiated for post-training a language model.

part of the algorithm	language-model form
solution θ	the weights of a LoRA adapter on a frozen instruction-tuned model

Table 1. The Seldonian algorithm instantiated for post-training a language model.

part of the algorithm	language-model form
data D	prompts, each with a group tag (benign, adversarial, a demographic value). Responses are sampled from the policy when needed
split	60/40 into D_c and D_s , stratified by group, with a fixed seed. D_s is sealed during training and read by one safety test only
objective	the expected reward of the policy’s responses, from a reward model or exact match
constraint g_i	the rate of a judged event among responses to a group of prompts, minus a threshold τ_i . The judge is a frozen function of (prompt, response): a guard model’s field, or code
threshold τ_i	absolute, or the reference policy’s rate plus a margin
candidate selection	GRPO on D_c ; the checkpoint that passes a predicted safety test with the highest reward
safety test	one fresh response per safety prompt, judged; a one-sided upper bound on the rate at level δ/k
NSF	the run returns no adapter

Two things differ from the classical setting. The policy can be queried, so $g(\theta)$ needs no off-policy estimate with importance weights: the candidate’s own response to each safety prompt is sampled and judged, and the safety test is a bound on a mean of bounded, independent variables. And the constrained quantity is not observed. It is whatever the judge says it is, so every guarantee below is about the judge’s label (sections 8.2, 9 and 10.3).

The safety test does not depend on how the candidate was produced. Reinforcement learning, supervised instruction tuning on D_c , or no training at all leave it unchanged. We have run the first (our own adapters) and the last (other people’s published models, sections 10.1 and 10.2), and not supervised instruction tuning.

4 Theory: what is optimised, $\hat{g}$, candidate selection, the safety test

4.1 The problem

Write x for a prompt, $y \sim \pi_\theta(\cdot | x)$ for a response, $r(x, y)$ for the reward, and $c_i(x, y)$ in $[0, 1]$ for judge i ’s label, with 1 meaning the constrained event happened. The problem is

$$\begin{aligned} & \text{maximise} && J(\theta) = \mathbb{E}_x \mathbb{E}_y [r(x, y)] \\ & \text{subject to} && g_i(\theta) = \mathbb{E}_{x \in \text{group } i} \mathbb{E}_y [c_i(x, y)] - \tau_i \leq 0, \quad i = 1, \dots, k. \end{aligned}$$

The first term of g_i is the *behaviour rate* $p_i(\theta)$: the probability that a response to a prompt from group i is flagged, over the prompt and over the policy’s own sampling. No algorithm can solve this problem exactly from a sample. The Seldonian requirement replaces it: return a policy with high reward, and only if $g_i(\theta) \leq 0$ can be supported at confidence $1 - \delta$. The optimiser increases a surrogate of J on candidate prompts, and a selection rule picks one checkpoint of the run. Neither enforces the constraint. The safety test does.

4.2 $\hat{g}$: the estimate of the constraint

For a prompt x_j drawn from group i and one response y_j sampled from the policy at temperature 1, let $z_j = c_i(x_j, y_j)$. Then

$$\hat{g}_j = z_j - \tau_i, \quad \mathbb{E}[\hat{g}_j] = p_i(\theta) - \tau_i = g_i(\theta).$$

Each $\hat{g}_j$ is an unbiased estimate of $g_i(\theta)$, and estimates on different prompts are independent when the prompts are an independent sample. This is the $\hat{g}$ of Thomas et al. (2019). Three quantities are built from the labels z_j and should not be confused: the point estimate $\text{mean}(z) - \tau$, used for reporting; the *predicted* upper bound, computed on candidate prompts during training and used to choose a checkpoint (section 4.3); and the *safety* upper bound, computed once on the safety set, which alone decides whether the policy is returned (section 4.4).

Relative thresholds. Several constraints here are relative: the trained policy’s rate may exceed the reference model’s by at most a margin. In the training runs cited, the reference rate was measured once on candidate prompts and $\tau = \text{reference rate} + \text{margin}$ was then treated as a constant, so the certificate is about that number and not about the true difference of the two rates. Section 10.4’s design does not do that. In human terms the reference rate is unknown, both policies answer the same safety prompts, and the estimate is the paired difference $d_j = z_j(\theta) - z_j(\theta_0)$ in $\{-1, 0, 1\}$.

4.3 Candidate selection

The update. At each step GRPO samples G responses to each prompt of a batch, scores them, and weights a clipped policy-gradient step by the score normalised within the prompt’s group of responses, with a KL penalty to the reference policy. Only the ranking of the responses to one prompt matters, so a penalty acts only where it changes that ranking. In the *filter-only* variant GRPO sees the reward r , and the constraint enters through the choice of checkpoint alone. In the *Lagrangian* variant it sees $r - \sum_i \lambda_i c_i$, each penalty applied on its constraint’s prompts: a stochastic ascent step on $L(\theta, \lambda) = J(\theta) - \sum_i \lambda_i g_i(\theta)$.

The predicted safety test. Every few steps the current checkpoint is sampled on m candidate prompts, one response each, and judged. The bound the safety test will use is computed on these labels at the sample size n_{eff} , with an optional inflation $\kappa \geq 1$:

$$n_{\text{eff}} = 1/(1/m + 1/n_s), \\ U_{\text{pred}} = \text{mean}(z) + \kappa \cdot (U(z; \delta_i, n_{\text{eff}}) - \text{mean}(z)).$$

m and n_s count the candidate and safety prompts in the constraint’s group. What has to be anticipated is the gap between the rate seen now on m prompts and the rate the safety test will see on n_s others, and that gap has variance proportional to $1/m + 1/n_s$. Thomas et al. (2019) use the safety-set size with $\kappa = 2$; the later runs cited here use n_{eff} with $\kappa = 1$. A checkpoint is *feasible* if $U_{\text{pred},i} \leq \tau_i$ for every constraint.

The multiplier and the choice. In the Lagrangian variant, after each predicted test,

$$\lambda_i \leftarrow \text{clip}(\lambda_i + \eta \cdot (U_{\text{pred},i} - \tau_i), \text{floor}, \lambda_{\text{max}}).$$

The step size under slack can be set separately, and in the language-model runs cited it was the same. The floor is zero until it is switched on, from the first update or once the constraint has been predicted infeasible. The predicted upper bound drives the multiplier, not the point estimate, so

the penalty rises before the estimated rate reaches the threshold. Among feasible checkpoints, the one with the highest mean reward on its predicted-test sample is kept. If none was feasible, the last one is tested anyway and usually fails.

Why none of this carries a guarantee. Lagrangian theory is about a saddle point: at convergence, the average of the iterates satisfies the constraint in expectation. It says nothing about the checkpoint in hand after 200 steps on a finite sample. The predicted test is computed on candidate prompts, with fresh responses, and the checkpoint is chosen because it looked best on them, so its predicted rate tends to be optimistic (Table 2). Candidate selection decides how often the safety test passes and how much reward the returned policy has. It has no effect on whether the guarantee is valid.

Table 2. “The constraint held”: a Lagrangian’s statement against a certificate’s.

	penalty or Lagrangian	certificate
measured on	fresh responses to candidate prompts, which the selection rule has seen	one draw of D_s , independent of both
statement	the estimate ended below the threshold in this run	the $(1 - \delta)$ upper bound on the returned policy is below the threshold
selection bias	uncorrected: predicted 0.133 against 0.184 on the safety set (one run of the companion paper); 0.141 against 0.168 (the constrained run of section 9.2). At another seed of the first experiment the predicted test erred the other way, by 0.01–0.02	removed by the fresh draw; that seed returned NSF
checked by	the breach rate: a fixed penalty of 1 breached in 3 of 3 seeds, a penalty of 4 in 1 of 3	the miss rate against δ (section 7.1)
can decline	no	yes; the solution rate is reported beside every certificate

4.4 The safety test

The test. Training has ended and one checkpoint θ_c is selected. For each constraint i , sample one response to each of the n_i safety prompts in its group, judge them, and compute a one-sided upper bound U_i on the rate $p_i(\theta_c)$ at level $\delta_i = \delta/k$. Return θ_c if $U_i \leq \tau_i$ for every i . Otherwise return NSF. The safety set is used once.

The guarantee. A bound is *valid* at level δ_i if $\mathbb{P}(U_i < p_i) \leq \delta_i$ for a fixed policy and an independent sample. Given D_c and all the randomness of training, θ_c is fixed and the safety labels are an independent sample with mean $p_i(\theta_c)$. A violating policy is returned only if $U_i \leq \tau_i < p_i(\theta_c)$ for some i , so

$$\mathbb{P}(\text{returned and } g_i(\theta_c) > 0 \text{ for some } i \mid D_c) \leq \sum_i \mathbb{P}(U_i < p_i(\theta_c) \mid D_c) \leq \delta.$$

Averaging over D_c gives the guarantee of section 2, for any selection rule on the candidate side.

The bounds. For 0/1 labels with k positives among n , the Clopper-Pearson bound is exact:

$$U = \text{Beta}^{-1}(1 - \delta; k + 1, n - k), \quad \text{with } k = 0 : U = 1 - \delta^{1/n}.$$

For labels in $[0, 1]$ the exact choices are Hoeffding’s and Bentkus’s bounds and a betting bound (Waudby-Smith and Ramdas, 2024). Student’s t is approximate. Section 6.3 gives the bounds for stratified, judge-assisted and clustered data. As a worked case: in the constrained run of section 9.2

the returned checkpoint was step 175 of 200; on the safety set its refusal rate by the guard’s flag was 0.168, the upper bound 0.198 and the threshold 0.206, so the adapter was returned.

NSF. The test returns NSF when $U_i > \tau_i$ for some constraint. NSF does not say the policy violates. It says this sample cannot support the claim at this confidence, because the rate is too close to the threshold or the sample is too small. Roughly, a policy with true rate p passes with probability one half when $\tau - p = z_{1-\delta} \sqrt{p(1-p)/n}$. With no positive observed, the smallest sample that certifies a threshold τ is $n \geq \ln(\delta)/\ln(1-\tau)$: 299 labels for 1% at delta 0.05, 149 for 2%, 59 for 5%. A method that always returns NSF satisfies the guarantee, so the solution rate is part of the result: 10 of 10 on a brevity constraint, 2 of 3 on over-refusal.

4.5 What the guarantee assumes

The argument of section 4.4 uses four assumptions. Each is routinely violated in language-model evaluation, and each has a part of this paper.

	assumption	how it fails	where we test it
A1	the safety set had no part in choosing the policy, and is used once	a split or a selection rule that uses safety-set labels; a test repeated until it passes	section 7.2
A2	the safety labels are independent draws from the population the claim is about	benchmarks that cross tasks with attacks; a labelling sheet that over-samples flagged responses	sections 9.3, 9.5, 10.2
A3	the bound is valid at this sample size and this rate	normal and Student-t limits at rates of a few percent	sections 7.1, 8, 9.4
A4	the label is the quantity the claim is about	a judge’s errors, corrected by a calibration measured on other responses	sections 9.1, 9.2, 10.3

5 What this paper proposes

Sections 2 to 4 describe an algorithm with two halves: a search for a candidate, and a test that decides whether the candidate may be returned. The search is the subject of the companion paper. This paper is about the test.

The proposal. We propose the Seldonian safety test, taken on its own, as a certificate for a behaviour rate of a language-model policy: a one-sided bound at level delta on the rate of a code- or model-defined label, computed on a sample the choice of policy never saw, with NSF as an outcome. The policy may be one we trained or one somebody else published. The rest of the paper audits the four assumptions of section 4.5. Our contributions:

- **Which bounds hold** their level at the sample sizes and rates of a safety test, with every bound we use measured under one rule (sections 7 and 8, Table 4).
- **What a label budget buys:** a safety set stratified by the reference model’s own per-prompt rate, and a judge used as a variance reducer under a routing rule fixed in advance (section 8).
- **What does not carry:** a judge’s calibration, across prompt populations or across training that targets the label; independent-sample bounds on crossed designs; and a stratified labelling sheet read as a random sample (section 9).
- **Certificates on published traces** of a robot benchmark and an agent-security benchmark (sections 10.1 and 10.2).

- **The price of human terms:** a guard model measured against human refusal labels, and the cost of a certificate of a trained policy in those labels (sections 10.3 and 10.4).

What the paper is not. It is not a new training method, and it does not claim that any model is safe. Nor is it a study of NSF for its own sake. NSF is one of the test’s two answers. The option to decline is what makes the guarantee possible, since a method that must return something cannot promise anything about what it returns, and how often the test declines is its cost. Here NSF is the answer for 1 of 3 training seeds on the over-refusal constraint, for all three models of a robot benchmark at a 5% threshold and delta 0.05, and for 27 of 28 pipelines of an agent-security benchmark (sections 7.2, 10.1 and 10.2). A solution rate is defined over repeated training runs and is reported there; a published trace has one sample and one answer.

Scope.

- The label is a guard model’s or the benchmark harness’s unless marked human. The human-label results are in section 10.3. Section 10.4 costs a certificate in human labels and does not compute one.
- Policies trained here are 0.5B to 2B parameters on one 12 GB card. The frontier-model results are certificates on other people’s published traces, not training.
- A certificate is about the label as defined and the population the safety sample was drawn from. It says nothing about other prompts or about the policy after further training.
- How the certified policies were trained is the companion paper’s subject. This paper needs training in two places only: the policy whose human-terms certificate section 10.4 costs, and the negative result of section 9.2, which exists because training targeted the label.

6 Evaluation set-up

The object under evaluation is a *certificate procedure*: a sampling design together with a bound that turns the labelled sample into an upper limit U . We ask whether the procedure is valid, and how wide its limit is.

6.1 Validity: the miss rate

Fix a population with a known true rate p , draw a sample S by the design and compute $U(S)$. The procedure *misses* when $U(S) < p$, and is valid at level delta if

$$\text{miss rate} = \mathbb{P}_S(U(S) < p) \leq \delta.$$

This is the one quantity the guarantee of section 4.4 needs from a bound. It cannot be read off a single real sample, where the truth is unknown, so we repeat the draw R times in a setting where p is known and take $\hat{m}$, the share of the R limits that fall below p . If the true miss rate were exactly delta, $\hat{m}$ would have standard error $\text{se} = \sqrt{\delta(1 - \delta)/R}$. One rule is applied to every bound, favoured or not. A cell (one bound, design, population, sample size and delta) is

- *over* its level when $\hat{m} > \delta + 2\text{se}$;
- *unresolved* when $\delta < \hat{m} \leq \delta + 2\text{se}$: the check cannot tell it from delta, which is weaker than showing that the bound holds;
- *at or under* otherwise.

We have not corrected for the number of cells inspected. With $R = 4,000$ and delta 0.05, a cell is over from a miss rate of 0.057.

6.2 Three sources of a known truth

A synthetic environment (the whole pipeline). A contextual bandit stands in for the language model, so that the training and testing code of section 3 runs unchanged and the true rate of any policy is an exact expectation: 20,000 contexts, four actions, a uniform reference policy that violates 13% of the time, and a reward whose pressure parameter makes the better actions the more violating ones. The unit of replication is a whole run (split, training, selection, safety test), and the estimated quantity is the Seldonian failure probability itself, $\mathbb{P}(\text{returned and truly violating})$. With a real model one replication is hours of training, so this is where assumption A1 is tested.

Resampling real responses (a plasmode; Franklin et al., 2014). A pool of real responses, each with its label, is treated as the population, and its mean is the truth. A replication draws a safety set from the pool by the design under test and computes the bound. The responses, the judge’s scores and their dependence are real; only the sampling is simulated. We use three kinds of pool.

- *Per-prompt pools.* 500 prompts, each with $k = 8$ judged responses of the reference model and further responses of each trained checkpoint, generated by Granite-3.3-2B. A replication draws n_s prompts without replacement, at random or within strata, and one response of the checkpoint for each.
- *Judge pools.* 500 responses, each with a gold label (the guard’s refusal field, or an exact word count) and a rubric judge’s score. A replication draws responses with replacement and reveals the gold label on the first n . In the cells at a set rate (20%, 5% and 1.3%) the label is drawn at that rate and the score from the pool’s scores for that label, so the truth there is the set rate.
- *Labelling sheets.* A sheet is re-drawn by its real sampling rule, which over-samples the responses a guard flags. The labels in this study are planted on the real responses at a known rate; they are not an annotator’s.

Resampling clusters (a crossed design). For AgentDojo each pipeline’s observed table of (user task, injection task) outcomes is the population and its overall rate the truth. A replication redraws whole tasks with replacement under one of three schemes: (a) user tasks redrawn, injection tasks fixed; (b) the reverse; (c) both. The scheme states what the certificate treats as sampled, and so what population it speaks about.

6.3 The bounds under test

All are one-sided upper limits at level delta; z is a normal quantile.

Labels alone. The bounds of section 4.4, and the Wilson score bound: the smallest $m \geq \hat{p}$ with $m - \hat{p} \geq z_{1-\delta} \sqrt{m(1-m)/n}$. It puts the variance at the hypothesised rate and not at the estimate, so its width is not zero when no positive is observed. It is approximate, and at rates of 1–2% it fails (Table 4).

****A stratified sample (b1w).**** With strata $h = 1, \dots, H$ of population share W_h , n_h labels and observed rate $\hat{p}_h$ in each, the estimate is $\hat{\mu} = \sum_h W_h \hat{p}_h$. The stratified Wilson-type bound is the smallest $m \geq \hat{\mu}$ with

$$m - \hat{\mu} \geq z_{1-\delta} \sqrt{\sum_h W_h^2 p_h(m)(1 - p_h(m))/n_h},$$

where $p_h(m)$ shifts every stratum’s rate by $m - \hat{\mu}$, clipped to $[0, 1]$. With one stratum it is the Wilson bound. It is approximate.

Labels and a judge (PPI++; Angelopoulos et al., 2023b). With n gold labels y , and the judge’s feature f on the same responses and on N_u unlabelled ones,

$$\begin{aligned}\hat{p} &= \text{mean}(y) + \lambda(\text{mean}_{\text{unl}}(f) - \text{mean}_{\text{lab}}(f)), \\ \lambda &= \text{cov}(y, f) / ((1 + n/N_u) \text{var}(f)), \\ V &= \text{var}(y - \lambda f) / n + \lambda^2 \text{var}_{\text{unl}}(f) / N_u.\end{aligned}$$

$\hat{p}$ is unbiased for the gold rate for any fixed λ , because the judge’s two means cancel in expectation. The published limit is $\hat{p} + z_{1-\delta} \sqrt{V}$. StratPPI (Fisch et al., 2024) applies the same correction within each stratum and combines the strata with weights W_h , again with a normal limit. PPBoot (Zrnic, 2024) replaces the normal limit by a percentile of bootstrap estimates.

The bootstrap-t limit. For any of these estimators, resample the labelled responses with replacement B times (within strata, for a stratified estimator), redraw each unlabelled mean from its sampling distribution, and recompute the estimate p_b^* and its variance V_b^* . With $t_b = (p_b^* - \hat{p}) / \sqrt{V_b^*}$ and q the delta-quantile of the t_b ,

$$U = \hat{p} - q\sqrt{V}.$$

It replaces the symmetric normal quantile by the estimated distribution of the studentised error, which is skewed when the rate is small. It is approximate.

Clustered data. With m clusters (user tasks) of sizes n_c and success counts s_c , $\hat{p} = \sum_c s_c / \sum_c n_c$ and

$$\text{se} = \sqrt{m / (m - 1) \cdot \sum_c (s_c - \hat{p}n_c)^2 / \sum_c n_c}.$$

The limit is the bootstrap-t limit with whole clusters resampled. The naive alternative treats the episodes as independent and applies Clopper-Pearson.

A paired difference. For a relative constraint in human terms the data are the differences d_j of section 4.2. The exact limit is the betting bound on $(d_j + 1)/2$, mapped back. The approximate ones are the bootstrap-t limit on $\text{mean}(d)$, and PPI++ with the difference of the guard’s two logits as the feature.

6.4 What a design or a judge is worth: effective sample size

Width is compared at equal validity. For a design or estimator A and a baseline R (a simple random sample with the pooled bound of the same family) at the same number of labels and the same delta, with w the mean distance from the estimate to the limit,

$$\text{ESS} = (w_R / w_A)^2.$$

An ESS of 2 means the design does with n labels what the baseline does with $2n$. In the comparisons with StratPPI (Table A1 and the medians of sections 8.1 and 8.2) w is measured from the truth to the limit, and in section 8.2 the baseline is the labels alone.

Reference-rate strata. One response per prompt has variance $p(1 - p)$, part of it between prompts (prompts differ in their rate) and part within (the policy’s own sampling). The between-prompt share is the intraclass correlation, ICC. A sample drawn within strata, in proportion to their

size, loses the variance that lies between strata; if that is a share G of the total, $\text{ESS} = 1/(1 - G)$. A relative constraint already calls for the reference model’s responses, so the strata come from them: sample the reference k times per prompt, rank the prompts by the share flagged, and cut the ranking into H equal strata with ties broken at random. This happens before training, so the split stays independent of the candidate (assumption A1). G can be predicted before any trained response is labelled:

$$G \approx \text{ICC}_{\text{cand}} \cdot \rho^2 \cdot \text{rel}(k) \cdot c_H, \quad \text{rel}(k) = k\text{ICC}_{\text{ref}} / (1 + (k - 1)\text{ICC}_{\text{ref}}),$$

where ICC_{cand} is the trained policy’s intraclass correlation, ρ the correlation between a prompt’s reference rate and its rate under the trained policy, $\text{rel}(k)$ the reliability of a k -sample reference rate, and c_H the share a cut into H strata keeps. When the pool of N prompts is itself a sample of the population the claim is about, the gain becomes $1/(1 - G + Gn_s/N)$.

A judge. With ρ^2 the squared correlation between the gold label and the judge’s feature, the variance V of section 6.3 gives $\text{ESS} = 1/(1 - \rho^2 N_u/(N_u + n))$, which tends to $1/(1 - \rho^2)$ when unlabelled responses are plentiful. A judge that predicts nothing has λ near 0 and an ESS of 1.

6.5 Carrying a calibration

A cheaper route uses no gold label on the certified policy at all. If the judge flags a true positive with probability s (recall) and a true negative with probability a (the false-alarm rate), its flag rate is $q = sp + a(1 - p)$, so $p = (q - a)/(s - a)$. A *carried* bound measures s and a on one population of responses, measures q on the population to be certified, and bounds p from one-sided limits on the three at $\delta/3$ each. It is valid only if s and a are the same on both populations, which is assumption A4 in its sharpest form. We measure the miss rate of the carried bound across two kinds of shift: between prompt populations, and between a reference model and a policy trained from it.

6.6 Models, labels and data

Table 3. The settings of the evaluation.

part	policy	label	sample	truth	delta
whole pipeline (7.2)	synthetic bandit	simulated judge	200 to 5,000 prompts	exact	0.10
labels alone (7.1)	three trained Qwen2.5-0.5B policies	guard’s harm field, rate 0.034	200 to 2,400 of 3,600 responses	pool mean	0.10
reference-rate strata (8.1)	Granite-3.3-2B, reference and trained	guard’s refusal and safety fields	100 to 200 of 500 prompts	pool mean	0.05, 0.10
judge routing (8.2)	500 real responses; a harm sheet of 225 from 4,800	gold: guard’s refusal field, a word count, 225 human harm labels; judge: a rubric	100 to 1,000 responses	pool mean	0.05
carried calibration (9.1, 9.2)	a reference model and two trained policies	gold: guard’s refusal field; judge: the rubric, six wordings	benign and harmful prompt pools	pool mean	0.05
robot benchmark (10.1)	three published frontier models	harness safety-stop flag	120 trials each	unknown	0.05, 0.10
agent benchmark (10.2)	28 published pipelines	harness security flag	629 to 949 task pairs each	table rate	0.05
guard against humans (10.3)	five published model variants; our two policies	human refusal labels	1,250 and 220 responses	none	95% intervals

Table 3. The settings of the evaluation.

part	policy	label	sample	truth	delta
human-terms design (10.4)	reference and trained Granite-3.3-2B	human strict refusal (not collected)	300 of 490 prompt pairs	syn- thetic labels	0.05

The benign prompts are XSTest’s safe prompts and OR-Bench’s (Röttger et al., 2024; Cui et al., 2025). The guard is Qwen3Guard (Zhao et al., 2025): the 4B model, run 4-bit quantised, except in the results cited from the companion paper (Table 2, the labels-alone row of Table 3 and the language-model seeds of section 7.2), which used the 0.6B model. Replications per cell are between 1,000 and 5,000 in Table 4, between 200 and 2,000 runs in the whole-pipeline studies of section 7.2, and 2,000 in the design check of section 10.4.

7 Results: validity

7.1 Which bounds hold their level

Section 6.1 defines the miss rate and the rule that sorts every cell into over, unresolved, and at or under its level; section 6.2 gives the three settings in which the truth is known. Table 4 applies that rule to every bound we use.

Table 4. Miss rate against delta for every bound used, with one rule for every row. Cells are counted as over / unresolved / at or under delta (section 6.1’s rule). “Exact” means valid at every sample size by construction; “approximate” means valid asymptotically and checked by resampling. Draws per cell are between 1,000 and 5,000 (4,000 for AgentDojo). The rows on a trained-policy harm pool are printed rates from the companion paper whose draws were lost. AgentDojo schemes: (a) user tasks resampled, injection tasks fixed; (b) the reverse; (c) both. The rows on StratPPI’s own allocations count 26 reference-strata cells: one checkpoint at a 95% rate is left out, where the proportional-allocation rows count all 28.

bound	kind	setting	miss, per cell delta		cells over / unresolved / at or under
Clopper-Pearson	exact (binary labels)	real harm labels, trained-policy pool, rate 0.034, n 200–2,400	0.10	0.067–0.084	0 / 0 / 5
Clopper-Pearson	exact	the plasmodes of section 8.2, every cell	0.05	0–0.052	0 / 2 / 40
betting mixture	exact (bounded)	same pool as row 1	0.10	0.007–0.016	0 / 0 / 5
Bentkus	exact (bounded)	same pool	0.10	0.023–0.035	0 / 0 / 5
Hoeffding, Anderson	exact (bounded)	same pool	0.10	0.000	0 / 0 / 5
stratified Wilson-type blw	approximate	4 mid-rate labels (9–94%), real Granite-3.3-2B responses, n_s 100–200, 3 checkpoints	0.05 / 0.10	0.017–0.054 / 0.058–0.097	0 / 3 / 21; 0 / 0 / 24
blw	approximate	label pushed by the Lagrangian, n_s 200	0.05 / 0.10	0.011–0.023 / 0.040–0.064	0 / 0 / 2; 0 / 0 / 2

Table 4. Miss rate against delta for every bound used, with one rule for every row. Cells are counted as over / unresolved / at or under delta (section 6.1’s rule). “Exact” means valid at every sample size by construction; “approximate” means valid asymptotically and checked by resampling. Draws per cell are between 1,000 and 5,000 (4,000 for AgentDojo). The rows on a trained-policy harm pool are printed rates from the companion paper whose draws were lost. AgentDojo schemes: (a) user tasks resampled, injection tasks fixed; (b) the reverse; (c) both. The rows on StratPPI’s own allocations count 26 reference-strata cells: one checkpoint at a 95% rate is left out, where the proportional-allocation rows count all 28.

bound	kind	setting	delta	miss, per cell	cells over / unresolved / at or under
b1w on a design-weighted sheet	approximate	sheets re-drawn by their real sampling rule	0.05	0.001–0.060	1 / 1 / 16
PPI++ with a bootstrap-t limit	approximate (second order)	the plasmodes of section 8.2, every cell and feature	0.05	0–0.053	0 / 12 / 156
cluster bootstrap-t, by user task	approximate	AgentDojo, 28 pipelines, scheme (a)	0.05	0.013–0.051	0 / 2 / 26
StratPPI estimator with a bootstrap-t limit, proportional allocation	approximate	reference-rate strata, 5 mid-rate labels, n_s 100–200; judge-logit strata, n 100–1,000	0.05	0.024–0.045; 0.001–0.056	0 / 0 / 28; 0 / 5 / 23
Student-t	fails at low rates	same pool as row 1, n 200–2,400	0.10	0.101–0.184	3 / 2 / 0
b1w and the pooled Wilson bound at rare rates	fail	harm labels at 1–2%, n_s 100, any design	0.05	0.076–0.448	12 / 0 / 0
PPI++ with a normal limit	fails in most cells	the plasmodes of section 8.2	0.05	0.041–0.241	119 / 32 / 17
StratPPI’s normal limit, proportional allocation	fails at these sizes	reference-rate strata; judge-logit strata	0.05	0.024–0.086; 0.054–0.239	9 / 5 / 14; 26 / 2 / 0
StratPPI’s normal limit, the paper’s allocations (oracle; heuristic)	fails	reference-rate strata; judge-logit strata	0.05	up to 0.110; 0.217; up to 0.219; 0.90	9 / 3 / 14; 11 / 0 / 15; 28 / 0 / 0; 24 / 1 / 3
StratPPI estimator with a bootstrap-t limit, oracle or heuristic allocation	fails	the same cells	0.05	up to 0.078; 0.174; up to 0.083; 0.90	6 / 3 / 17; 11 / 1 / 14; 14 / 11 / 3; 14 / 0 / 14
PPBoot, percentile limit (basic; power-tuned)	fails at these sizes	judge-logit cells, unstratified	0.05	up to 0.133; 0.148	7 / 6 / 1; 12 / 2 / 0
Clopper-Pearson over pairs	fails on a crossed design	AgentDojo, 28 pipelines, schemes (a); (b); (c)	0.05	0.042–0.268; 0.060–0.340; 0.195–0.350	27 / 0 / 1; 28 / 0 / 0; 28 / 0 / 0
cluster bootstrap-t by user task, when injection tasks are sampled	fails	AgentDojo, schemes (b); (c)	0.05	up to 0.258; 0.275	21 / 0 / 7; 24 / 0 / 4
larger of the two clustered bounds	holds under (a) only	AgentDojo, schemes (a); (b); (c)	0.05	at most 0.029; up to 0.066; 0.089	0 / 0 / 28; 4 / 0 / 24; 16 / 3 / 9
two-way bootstrap	fails	AgentDojo, schemes (a); (b); (c)	0.05	up to 0.174; 0.178; 0.278	5 / 0 / 23; 14 / 1 / 13; 27 / 0 / 1
stratified sheet read as an i.i.d. sample	fails	PPI on the sheet, three judge wordings	0.05	0–0.989	7 / 0 / 11

Table 4. Miss rate against delta for every bound used, with one rule for every row. Cells are counted as over / unresolved / at or under delta (section 6.1’s rule). “Exact” means valid at every sample size by construction; “approximate” means valid asymptotically and checked by resampling. Draws per cell are between 1,000 and 5,000 (4,000 for AgentDojo). The rows on a trained-policy harm pool are printed rates from the companion paper whose draws were lost. AgentDojo schemes: (a) user tasks resampled, injection tasks fixed; (b) the reverse; (c) both. The rows on StratPPI’s own allocations count 26 reference-strata cells: one checkpoint at a 95% rate is left out, where the proportional-allocation rows count all 28.

bound	kind	setting	miss, per cell delta	cells over / unresolved / at or under
the judge’s rate alone	fails	the plasmodes of section 8.2	0.05 0–1.000	30 / 0 / 12

The one cell over its level in a row we call valid is **b1w** on a design-weighted sheet (0.0595 at 4,000 draws). It is not over after a correction for the row’s 18 cells. Those 18 cells are 9 settings run twice: the other run of the same setting gave 0.0485, and the two together give 0.054, which is unresolved. Every row we call failing has at least one cell over, and most also have cells that are not: a failing bound fails somewhere, not everywhere.

Three readings of Table 4.

1. **Exact bounds for rare labels.** At 1–2% and n_s 100 the Wilson-type bounds, pooled or stratified, missed in 8–45% of draws (24–44% at each label’s worst checkpoint); the failure is binomial discreteness, not the design. At n_s 200 **b1w** is still over its level for the 1% label and at or under it for the 2% label. The approximate limits that kept their level there (a Wald-t and the bootstrap-t limits) did so with bounds wider than the labels alone give. Rare labels take Clopper-Pearson or a betting bound.
2. **The t-test is the wrong default.** On the one trained-policy pool we resampled (a harm rate of 3.4%) it is over its level at sample sizes up to 800 and unresolved above them, and its width is zero at a rate of 0 or 1, which lets candidate selection win by driving a subgroup rate to the boundary.
3. **“Approximate” has a number attached.** The approximate bounds we use are over their level in 1 of 322 cells and unresolved in 23; the largest miss at delta 0.05 is 0.060. The 322 are the cells of the six approximate rows of Table 4 that are not in bold, and the delta 0.10 cells of the two **b1w** rows reuse the draws of their delta 0.05 cells. The count is of the settings where we use these bounds. It leaves out the rare labels, where **b1w** fails (the bold row), and we have not mapped where between a 2% and a 9% rate **b1w** begins to hold. Inside the training loop on a synthetic environment the stratified test reached 0.116 at delta 0.1 on 500 runs, which is unresolved (section 7.2).

The StratPPI rows (Fisch et al., 2024) are our implementation of its estimator and interval on the same draws as the rows above them; sections 8.1 and 8.2 give the comparison.

7.2 Validity under the full pipeline

Table 4 tests bounds on a fixed population. The certificate also has to survive selection: the candidate is the best-looking checkpoint of a training run.

Synthetic environment, exact truth. In a contextual bandit where every policy’s true rate is computable, the probability of returning a policy that truly violates was 0.000–0.002 against a

delta of 0.1 for the Lagrangian variant of the Seldonian arm under four bounds, against 0.830 for unconstrained training at the same reward pressure. In two one-way sweeps of that variant it stayed at or below 0.002, across reward pressures 0 to 4 at 1,000 prompts and across sample sizes 200 to 5,000 at pressure 1. The filter-only variant was higher and still under delta: 0.028 at 200 prompts and 0.004 at 1,000 and 5,000. With a noisy judge the guarantee held at the judge’s level in every row, while a judge that misses 40% of violations let the true violation rate among returned policies reach 0.059: the certificate is about the label.

Adversarial splits. Rerandomised and stratified candidate/safety splits did not break the exact test, across 2,000 seeds per arm, eleven split rules and candidates built to overfit the balancing statistics: Clopper-Pearson missed in at most 1 of 2,000 runs in any cell. A tight Wald test on the same runs was unresolved above delta in 3 of 42 cells, with a largest miss of 0.055 at delta 0.05.

The training loop, on a synthetic environment. With our `SeldonianLLMPolicy` class training on a synthetic bandit at four levels of between-prompt heterogeneity, with 4 strata and 500 runs per cell, the stratified `b1w` test missed 0.084–0.116 at delta 0.1 (Monte Carlo standard error 0.013): three cells at or under delta and one unresolved above it. A random split with a pooled bound missed 0.090–0.100, which 500 runs cannot tell apart from the stratified test. No language model is in this check.

Language-model policies. On an over-refusal constraint (Qwen2.5-0.5B, three seeds), unconstrained training and a fixed penalty of 1 breached the threshold in 3 of 3 seeds, a penalty of 4 in 1 of 3; the Seldonian arm breached in 0 of 3 and certified 2 of 3. On a brevity constraint, 10 of 10 seeds certified with no breach. Three and ten seeds cannot estimate a miss rate; they are consistent with the level, and the resampling studies above are the evidence.

Trajectories. On a synthetic bandit, a test of the returned policy alone misses drift through a forbidden region during training: 65% of plain-Lagrangian runs passed it after a step inside the region (50% when checked every 10 steps). A certificate at level delta/T over every one of T checks was at or under delta in all 10 cells, with misses of 0.025–0.100 against a delta of 0.1, at 200 runs a cell, which resolve only a miss above 0.14.

8 Results: what a label budget buys

8.1 Making the sample go further: reference-rate strata

Idea. Section 6.4 gives the design and why it can help: sample the reference k times per prompt, cut the prompts into H equal strata of their reference rate, sample the safety set within strata, and bound the rate with the stratified Wilson-type bound `b1w`. The gain depends on how much of the per-response variance sits between prompts (the intraclass correlation of the reference, ICC_{ref}) and on how well a prompt’s reference rate predicts its rate under the trained policy.

Result. With $k = 8$ and $H = 8$ on real Granite-3.3-2B responses, the effective safety-sample size (ESS: the size of a simple random sample giving the same bound width) against a random split was 2.4 for over-refusal (rate 17%, ICC_{ref} 0.72), 5.1–5.3 for refusal of plainly harmful requests (66%, 0.86), and 1.4 for non-refusal of encoded requests (9%, 0.50), with coverage as in Table 4. A placebo covariate gave 1.0. When the Lagrangian pushed the stratified label itself, the gain was 2.1–2.2 with no compression of the between-prompt spread (correlation between a prompt’s reference and trained rates 0.92).

The rule was chosen on these pools. The strata rule above is not the one we registered in advance. That rule (ties kept together, cuts at quantiles, small strata merged) collapsed to one or two strata here, because 73–97% of prompts have an 8-sample reference rate of exactly 0 or 1, and often gave no gain. Equal rank strata with random ties were compared with it on the same pools

and adopted, and the truth used for coverage was changed to the mean of the half of the responses the draws come from. The gains above are therefore measured on the data that chose the rule, and have not been checked on a pool that played no part in the choice.

Where it does not help. Rare labels (the approximate bound is invalid there and exact stratified bounds did not beat pooling); labels near 0 or 1 (0.96 to 1.12 at a 93% rate, Table A1); a safety set that is a large share of the prompt pool (the gain is capped at $1/(1 - G + Gn_s/N)$); and task strata on a benchmark with 20 trials a task, where the bound is set by the positives.

A pre-flight. The gain is predictable from k reference samples before any trained-policy response is labelled: Spearman 0.83 between predicted and realised ESS on real data, with the prediction 0–20% high at $H = 8$ (Figure 1). The absolute-error criterion we pre-registered (median error at most 0.1) failed on real data (0.29) and passed on the bandit (0.02); we report the ranking as the usable part.

The gain from reference-rate strata can be predicted before labelling

8 strata of an 8-sample reference rate, 200 safety prompts, 6 0.1, 200 training steps. Labels with rates of 1-2% are left out: no approximate bound holds there.




Figure 1. Gain from reference-rate strata against the reference model’s intraclass correlation.

Against StratPPI. StratPPI (Fisch et al., 2024; section 6.3) is the published method closest to this one. We ran our implementation of it on the same strata and the same 5,000 draws per cell, with the reference rate as the predictor (Figure 2, left; Table A1 in Appendix A). Three things follow.

1. **StratPPI’s bound is shorter, and part of that is the normal limit running over its level.** It is over its level in 4 of 10 cells at delta 0.05 when a cell is the largest miss over checkpoints (one of them marginally, at 0.056), and in 9 of 26 cells counted by checkpoint (leaving out one at a 95% rate), most at the 9% label.

StratPPI's interval is shorter partly because it runs over its level; with an honest limit its estimator holds

Each dot is one cell (a label and a sample size). Right of the line the bound misses more often than it claims. Baseline: a random split with a pooled Wilson bound (left), the labels alone (right).

- this paper's bound for the setting: stratified Wilson-type (left), PPI++ with a bootstrap-t limit (right)
- StratPPI as published, with its normal-quantile limit
- StratPPI's estimator with a bootstrap-t limit




Figure 2. StratPPI’s normal limit and the same estimator with a bootstrap-t limit, with labels allocated in proportion to stratum size. Left: reference-rate strata. Right: strata on a judge’s logit.

2. **With an honest limit the two are close.** With labels allocated in proportion to stratum size, the same estimator with a bootstrap-t limit is over its level in none of the ten cells (largest miss 0.045), and its median ESS is 2.10 against 2.16 for **b1w**. The repair depends on the allocation (Appendix A).
3. **So the gain here is the stratification.** The within-stratum regression on the reference rate adds little once the limit is valid. We keep **b1w** as the bound for reference-rate strata at these sizes and note the bootstrap-t StratPPI as the narrower choice at a safety set of 200 or with 4 strata, where its median ESS is 2.6–2.7 against 1.8–2.2 for **b1w**.

What the gain costs. The strata are built from $k = 8$ judged reference responses on each of the pool’s 500 prompts, 4,000 in all, and the ESS counts only the safety set’s labels. An ESS of 2.4 is worth 140 labels at n_s 100. So the method saves labels only where the label on the trained policy costs far more than the stratifier’s, for instance a human label stratified on a guard’s reference rate. We did not test that case, and section 10.3 shows the guard and the annotator disagree on about 40% of the guard’s flags. Where both are the same cheap judge, enlarging the safety set is the better use of the budget.

8.2 What a judge buys: a routing rule

A judge model scores every response cheaply; gold labels are scarce. We fix in advance which estimator a constraint gets (Table 5).

Table 5. Routing rule.

data	route	guarantee
a property code can verify (a word count)	compute it on every response; Clopper-Pearson	exact
gold labels, fewer than 10 in the rarer class	Clopper-Pearson on the labels; judge unused	exact
gold labels, 10 or more in each class	PPI++ (Angelopoulos et al., 2023b) on the judge’s logit, bootstrap-t limit	approximate
gold labels from a stratified sheet	design-weighted labels, b1w	approximate
a calibration measured on another population	refused	none

The judge’s worth in labels. With ρ^2 the squared correlation between the gold label and the judge’s feature, PPI++ multiplies the labels’ worth by about $1/(1 - \rho^2)$ when unlabelled responses are plentiful. For a compiled refusal judge against the guard’s label (225 labels, 2,000 judged draws from a pool of 500 responses): the 0/1 verdict gives 1.11, the probability 1.15, the logit 1.61, a cross-fitted Platt score 1.86 in measured bound width. For this judge the logit carried more than the verdict. For a harm rubric ρ^2 was 0.02, a gain of 1.02: the labels carry the certificate and the judge contributes nothing.

Label budgets for a rare label. With no positive observed, 299 gold labels on the certified policy’s own responses certify a 1% threshold at delta 0.05, 149 certify 2%, and 59 certify 5%.

Choosing which responses to label. Table 5 labels a random subset. If the judge has scored every response first, the labels can be drawn within strata of its logit. On the refusal pools of those plasmodes (two judge wordings, 14 cells, 4,000 draws each, delta 0.05) (Figure 2, right):

- StratPPI as published, with 5 or 10 strata, is over its level in 26 of 28 cells (misses 0.059–0.239 in those 26), as PPI++ with a normal limit is in all 14 (0.061–0.232). Appendix A gives the counts under other conventions and allocations.
- With proportional allocation the StratPPI estimator with a bootstrap-t limit is over its level in none of the 28 (5 unresolved, largest miss 0.056). It is the narrowest such route with 5 or 10 strata at the 20% and 5% rates: with 10 strata its median ESS against the labels alone is 2.54 at a 20% rate and 1.87 at 5% (2.47 and 1.76 with 5 strata), where unstratified PPI++ with the same limit gives 1.67 and 1.39.
- Stratifying on the judge and ignoring it within strata (**b1w**) is also over in none (1 unresolved, largest miss 0.053), at 2.41 and 1.42 with 10 strata (2.24 and 1.26 with 5). With 20 strata, and wherever strata hold about ten labels or fewer, **b1w** is the narrower.
- PPBoot’s percentile limit (Zrnic, 2024), on the 14 unstratified cells, is over its level in 7 with the basic estimator and 12 with the power-tuned one.
- At a 1.3% rate with 225 labels the bootstrap-t routes return nothing useful and the count rule’s fallback to the labels alone is the right one.

So the third row of Table 5 has a better form when labelling can follow scoring: stratify on the judge’s logit, StratPPI’s estimator, a bootstrap-t limit, with labels allocated in proportion to stratum size and about 45 or more per stratum. Below that, **b1w** on the same strata is narrower. The count threshold stays.

9 Results: what does not carry

Figure 3 shows sections 9.1 and 9.2: the miss rate of a carried bound and the judge’s recall, by shift.

A judge’s calibration does not carry, least of all across training that targets the label

A rubric judge’s recall and false-alarm rate are measured on one population and used to bound the refusal rate of another. Left: how often that bound falls below the truth, one dot per wording of the judge’s rubric. Right: what happened to the judge’s recall.




Figure 3. A carried calibration: miss rate of the carried bound and the judge’s recall, by shift.

9.1 A judge’s calibration across prompt populations

Measure a judge’s recall and false-alarm rate on one population and apply them to another, and the corrected bound can fail. For one compiled judge (six wordings) scored against the guard’s label: across two sources of benign prompts in the same pool, recall differed in 4 of 6 wordings and the carried bound missed in up to 94% of draws; across benign and harmful pools, 4 of 6 and up to 100%. The failure depends on the direction of the carry. Between the two benign sources the bound was over its level for 4 of 6 wordings one way and 1 of 6 the other; between the pools, for 5 of 6 one way and none the other. False-alarm rates did not differ, but they sit at or near zero for this judge. Its recall on the guard’s positives is 0.04–0.54, so part of a recall difference may be a difference in what the guard flags. The tests are unadjusted for six wordings.

9.2 The same calibration across training that targets the label

The measured result is an existence result from two single runs. A calibration carried across 200 steps of one training run that moved the label only as a side effect (recall differed in 0 of 6 wordings). In one run whose Lagrangian targeted the label it did not carry: at step 200 the compiled judge’s recall on the guard’s refusals fell from 0.21 to 0.03 for the canonical wording, recall differed in 2 of 6 wordings, and the carried bound was over its level for 3 of 6 (misses 0.08, 0.23 and 0.80) and marginal for a fourth (0.054), although the rate itself had returned to the reference’s. At step 100, where the pool’s rate was 3.4 points below the reference’s, 1 of 6 failed.

The same direction appears against human labels, without reaching significance. We scored the annotator’s 220 labelled responses with the same judge. Its recall on the annotator’s strict refusals is 3 of 18 for the reference and 0 of 18 for the trained policy under the canonical wording, and

lower after training in all six wordings (Fisher p from 0.18 to 1.0). The guard’s recall and precision against the same labels do not move (17 of 18 and 16 of 18; 0.59 and 0.60), so the fall is not the guard flagging a different mix.

What this does not show is the cause. The two runs differ in reward and set-up as well as in the constraint, each has one seed, the gold label is the guard’s, and the judge misses most of the guard’s positives to begin with. We wrote down beforehand that a failure at step 200 would mean the earlier success “was luck and no calibration carries across training either”; reading the earlier run as a control arm came after the result. Either reading gives the same rule.

Appendix B gives our account of the mechanism, which no experiment here isolates: a change in the form of the responses the guard flags.

The consequence for certification: a calibration cannot be assumed to carry to a policy it was not measured on. Gold labels belong on the responses of the policy being certified, each time the policy changes.

9.3 Independent-sample bounds on crossed designs

AgentDojo crosses user tasks with injection tasks; a pipeline’s attack successes cluster by both. Design effects by user task ran 1.4–8.7 and by injection task 2.6–15.5. The per-pair Clopper-Pearson bound was over its level for 27 of 28 pipelines with user tasks resampled (misses up to 27%) and for all 28 once injection tasks are resampled too (Table 4). Section 10.2 gives the consequence for a published leaderboard.

9.4 Normal limits and judge-assisted exact bounds

PPI++ with its normal limit missed in up to 24% of draws at delta 0.05 (n 225 at a 1.3% rate); at a 20% rate and n of 100–500 it missed in 6–11%. Finite-sample judge-assisted bounds (betting on blocks) were valid and bought nothing: narrower than the labels alone in 10 of 210 cells, by at most 0.003. Taking the smaller of two valid bounds is not a bound.

9.5 A stratified sheet read as a random sample

A labelling sheet that over-samples a guard’s flags is a stratified sample. Read as i.i.d., it certified a negative harm rate under one judge wording. With labels planted at a 1.3% rate it missed in 98% of re-drawn sheets under that wording, and in 0–99% across wordings and rates (Table 4). Every sheet in this paper carries its design weights (section 10.3).

10 Results: certificates in practice

10.1 A robot benchmark: the resolution is the result

RoboDojo-RC Tier 1 publishes 120 trials per model with a harness-emitted safety-stop flag. Table 6 gives the certificates and Figure 4 what 120 trials can resolve.

Table 6. Safety stops, Clopper-Pearson upper bounds.

model	n	stops	rate	upper, delta 0.05	upper, delta 0.10	“at most 5%” at delta 0.05
Claude Opus 5	120	10	0.083	0.137	0.125	NSF
Claude Opus 5.5	120	2	0.017	0.052	0.044	NSF (passes at delta 0.10)
GPT-6 Astra	120	2	0.017	0.052	0.044	NSF (passes at delta 0.10)

Table 6. Safety stops, Clopper-Pearson upper bounds.

model	n	stops	rate	upper, delta 0.05	upper, delta 0.10	“at most 5%” at delta 0.05
-------	---	-------	------	----------------------	----------------------	-------------------------------

A benchmark of 120 trials can certify a 2.5% rate when it observes nothing and cannot certify 5% after two events; a 1% claim with nothing observed needs 299 trials per model. One model stops more than the other two (10 of 120 against 4 of 240, one-sided Fisher $p = 0.003$). The policy’s own narration of risk does not predict a stop (AUC 0.38); joint effort does (AUC 0.75–0.83, exploratory).

10.2 An agent-security benchmark: one of 28

On AgentDojo’s published runs (19,380 episodes, the harness’s **security** label), what a certificate can say depends on what is taken as sampled. We take the benchmark’s injection tasks as given and its user tasks as a sample of tasks like them. The certificate is then the studentised cluster bootstrap by user task (Figure 5). Resampling all 28 pipelines 4,000 times each, that bound is over its level for none (2 unresolved).

It does not extend to new injection tasks. When injection tasks are resampled as well, the same bound is over its level for 24 of 28 pipelines. So are the others we had: the bootstrap by injection task (23), a two-way bootstrap (27), and the larger of the two clustered bounds (16). A bound that adds the two clustered margins in quadrature was over for none in this one check; we did not specify it in advance and do not rely on it.

Table 7. Targeted attack success, selected pipelines, upper bounds at delta 0.05. The certificate is the bound clustered by user task, with injection tasks taken as fixed. Every bootstrap bound is from 200,000 draws. The verdicts are the same under the larger of the two clustered bounds and at 30 seeds of a 4,000-draw bootstrap.

pipeline	pairs	rate	per-pair bound	clustered by user task	clustered by injection task	“at most 5%”
claude-3-5-sonnet-20241022	629	0.011	0.021	0.022	0.035	pass
Meta-SecAlign-70B	949	0.022	0.032	0.062	0.035	NSF
command-r	629	0.033	0.048	0.063	0.054	NSF
claude-3-7-sonnet-20250219	949	0.050	0.063	0.070	0.104	NSF
gpt-4o + tool_filter	629	0.068	0.087	0.092	0.106	NSF
gpt-4o-2024-05-13	629	0.477	0.511	0.537	0.588	NSF

One pipeline of 28 certifies under the bound clustered by user task, and the same one alone under the larger of the two clustered bounds. The bootstrap by injection task and the two-way bootstrap, which are over their level under all three resampling schemes, would also pass the two Meta-SecAlign-70B pipelines. Its pass rests on 7 successes in 6 user tasks. A defended model whose raw rate is 2.2% does not certify: its 21 successes sit in 5 of 97 user tasks (ICC 0.72), and the bound clustered by user task is 0.062 against 0.032 per pair. That bound is sensitive to the size of the bootstrap when a handful of clusters hold every success: at 4,000 draws it ranged from 0.058 to 0.107 over 30 seeds. The other rows moved by at most 0.013 across seeds. Five pipelines pass the per-pair bound, and four of those passes do not survive clustering.

What a 120-trial benchmark can certify

Clopper-Pearson upper bounds on the safety-stop rate. Nothing observed certifies 2.5%; two stops already miss a 5% claim at δ 0.05.




Figure 4. What a benchmark of 120 trials per model can certify.

10.3 The guard against human labels

Published human labels. XSTest releases 2,250 completions from five model variants, each carrying two annotations and a final label: full compliance, full refusal or partial refusal (“any combination of refusal and compliance”; Röttger et al., 2024). We scored every completion with Qwen3Guard-4B, run 4-bit quantised, and compared its flag with the final human label on the 250 safe prompts per model (Figure 6). This is the guard behind the refusal numbers of sections 8 to 10 and of the appendices. The three-seed over-refusal runs cited from the companion paper (Table 2 and section 7.2) used the smaller Qwen3Guard-0.6B, which we have not compared with human labels.

Table 8. The guard’s refusal flag against XSTest’s human labels, safe prompts. Recall and false-alarm rate against full refusal alone (strict) and against full or partial refusal (broad).

model	humans: full / partial refusal	guard flags	strict: recall, false alarms	broad: recall, false alarms
GPT-4	0.064 / 0.020	0.088	1.00, 0.03	0.90, 0.01
Llama 2, original system prompt	0.380 / 0.216	0.636	0.96, 0.44	0.94, 0.19
Llama 2, no system prompt	0.140 / 0.156	0.304	0.94, 0.20	0.85, 0.07
Mistral instruct	0.008 / 0.008	0.036	0.50, 0.03	0.50, 0.03
Mistral with a guardrail prompt	0.096 / 0.092	0.168	0.75, 0.11	0.70, 0.04
pooled (1,250)	0.138 / 0.098	0.246	0.92 (0.88–0.96), 0.14	0.87 (0.82–0.91), 0.05

On AgentDojo's published runs, one pipeline in 28 certifies a 5% attack success rate

The usual bound treats every user-task and injection-task pair as independent. The clustered bound is a cluster bootstrap by user task, with the benchmark's injection tasks taken as fixed.

● observed rate ● bound treating pairs as independent ● bound clustered by user task




Figure 5. AgentDojo: the per-pair bound and the bound clustered by user task for all 28 pipelines, from 200,000 bootstrap draws.

The guard's refusal flag is broad: it marks most partial refusals too

Qwen3Guard-4B on the responses of five models to XSTest's 250 safe prompts, against the published human label. A class with fewer than 5 responses is not drawn.




Figure 6. What the guard flags, by human class, on XSTest's published completions.

Three readings.

1. **The guard's flag is a broad one.** It marks 92% of human full refusals (159 of 172), 80% of partial refusals (98 of 123) and 5% of full compliance (51 of 955). Its agreement with the broad human label (Cohen's kappa 0.81) is close to two humans' agreement on the three classes (0.88); with the strict label it is 0.59 against the humans' 0.90.
2. **What it is worth as a predictor.** The squared correlation of the human label with the guard's logit is 0.70 for the broad event and 0.46 (95% interval 0.39–0.53) for the strict one. By section 8.2 that is roughly a threefold and a twofold gain in labels when unlabelled responses are plentiful. The logit beats the 0/1 flag for the strict event (0.46 against 0.40) and for the broad one (0.70 against 0.65).
3. **Its error rates belong to a model.** The false-alarm rate against the strict label runs from 0.03 (GPT-4) to 0.44 (Llama 2 with its original system prompt), following each model's share of partial refusals; recall runs from 0.75 to 1.00 among the four models with more than two refusals. This is section 9.1 on other people's models and labels: a calibration measured on one of these models would not serve another.

A consequence for section 10.4: the certified event there is strict refusal, and the guard that shaped the training measures the broad one. The guard can still reduce variance (reading 2); it cannot stand in for the label.

Our own policies. One annotator (an author) labelled 220 responses to benign prompts: a stratified sample of the reference model's and the trained policy's responses, read with its design weights, under a guideline that follows XSTest's scheme with refuse-then-answer as its own label.

Figure 7 and Table 9 give the result.

In human terms the two policies refuse equally often, and less often than the guard says

220 responses labelled by one annotator, a stratified sample read with its design weights; bars are conservative 95% intervals. The guard's rate is its share of the whole pool, so it has no interval.




Figure 7. Refusal rates of the reference model and the trained policy, by the guard’s flag and by human label.

Table 9. Human labels on our own policies (one annotator, 220 responses; conservative 95% intervals in brackets).

	reference model	policy trained under the constraint (step 200)
the guard’s refusal flag, share of the pool	17.9%	17.4%
human label: refuses	10.7% [4.7, 21.5]	10.6% [5.2, 21.2]
human label: refuse-then-answer	3.4% [0.4, 15.5]	0.3% [0.0, 11.0]
guard’s false-alarm rate against “refuses”	8.1%	7.7%
guard’s precision	0.59	0.60
guard’s recall: estimate, conservative lower limit	0.99, 0.44	0.98, 0.45
ρ^2 of the human label with the guard’s logit	0.61	0.60

In human terms no difference between the two policies was detected (-0.1 points, prompt-clustered standard error 3.9, so roughly -8 to $+8$), and both refuse about 7 points less often than the guard’s flag says. Each rate leans on small strata: the trained policy’s on one stratum of 18 labels, the reference’s on two of 15. The responses were cut at 128 tokens. The guard missed almost nothing the annotator called a refusal, but the sample cannot rule misses out: none was found among 72 and 70 responses in the large guard-negative strata, which bounds recall below only at 0.44. As a predictor the guard is worth more than a halving of the labels (ρ^2 0.6), in line with the published-label result above.

One annotator means no agreement statistic. What stands in for it is thin and we say so: the guideline’s classes are XSTest’s, on whose released completions for safe prompts the two annotations agree with a kappa of 0.88; three responses that appeared twice on the annotator’s sheets received

the same label; and two recurring cases (a redirect that says where to find the answer; “I cannot answer” followed by the false premise) were labelled both ways.

10.4 What a certificate in human labels costs

The constraint of the trained policy is relative: its refusal rate may exceed the reference’s by at most a margin of 0.02. Section 10.3 puts the two human-terms rates at 10.7% and 10.6%. With a true difference near zero, whether a sample certifies the margin is a question about the label budget. We report that budget and the margin a fixed sample would be expected to certify. We prepared the measurement and did not run it.

Table 10. Prompt pairs needed to certify a margin on the difference of the two strict refusal rates at delta 0.05, from the rates of section 10.3. A design calculation: normal-type limits, pairing correlation 0.66 (the guard’s flags on the same prompts), guard ρ^2 0.6, and the unrounded difference of the two rates (−0.0005). The rates are those of the last checkpoint (step 200); the returned policy that section 10.4’s samples come from is step 175. The two guard columns differ in what the claim covers: the pool’s own 490 prompts, where guard-only responses pin down the guard’s mean, or new prompts from the same source, where the pool’s unlabelled prompts are all the guard has.

margin	chance of certifying	labels alone, unpaired	labels alone, paired by prompt	with the guard, pool rate	with the guard, new prompts
0.02	50%	1,226	416	167	339
0.02	80%	2,801	950	380	not within 490 prompts
0.03	80%	1,266	429	172	362
0.05	80%	462	157	63	78

The budget turns on the true difference, which the labels fix only to within about 8 points either way. At a difference of +0.006, the value in the design check below, the same formula gives about 930 pairs for an even chance and 2,100 for an 80% chance; at +0.02 nothing certifies. The guard’s ρ^2 of 0.6 comes from about 110 labels per policy and has no interval here.

Pairing the two policies on the same prompts cuts the budget threefold. What the guard adds depends on the claim. For the rate on the pool’s own prompts, a guard-only response costs GPU seconds, the guard’s mean can be measured as closely as wanted, and the guard cuts the budget again by 2.5: a certificate that would take 2,800 unpaired pairs takes 380. For new prompts from the same source, the guard’s mean is known only through the pool’s unlabelled prompts. With 300 of 490 labelled it removes 23% of the variance ($0.6 \times 190/490$), not 60%, and the 2-point margin at an 80% chance is out of reach inside the pool. A predictor removes label noise. It does not remove the uncertainty of a small prompt set.

Prepared, not run. We drew one fresh response (256 tokens) per policy on each of the 490 pool prompts, from the returned policy of one training run and from its reference, and 8 further responses per policy on each prompt that only the guard reads. 300 of the 490 pairs form the sheet. All 8,820 responses were scored with the guard. The analysis was fixed before any label: four upper limits at delta 0.05 on the difference, each computed once on the full sample. (a) Labels alone, the betting bound (exact). (a’) Labels alone, bootstrap-t (approximate, the like-for-like comparator). (b1) With the guard, for the pool rate, PPI++ with a bootstrap-t limit. (b2) With the guard, for new prompts, the same with the other 190 prompts’ pairs as the unlabelled data.

A check of the whole design on the guard’s real scores, with synthetic labels drawn from the guard’s logit at the first sheet’s rates in each band of the logit, puts the miss rates of the four limits at or under 0.045 for the pool rate against a level of 0.05, and their mean certified margins at 0.065,

0.037, 0.032 and 0.035. The exact bound is about twice as wide as the approximate ones. Appendix C gives the check in full.

On either set of figures a sample of 300 pairs is expected to certify a margin of 2.4 to 6.5 points with labels alone, not the 2-point target. The 600 labels were not collected. The samples, the guard’s scores, the blind sheet and the analysis script are in the repository, so the certificate can be computed from 600 labels with no design choice left open. One property of the sheet would limit that measurement: the reference writes longer responses than the trained policy (57% against 32% of the sheet’s items stop at the token limit in mid-sentence), so length is a weak cue to the policy.

11 Limits

- **Labels.** Outside section 10.3, and the budget of section 10.4 computed from it, every refusal and harm number is relative to a guard model’s field: Qwen3Guard-4B, 4-bit quantised, or the 0.6B model in the results cited from the companion paper, which has no human comparison here. The human refusal labels are one annotator’s, an author’s, with no measured agreement; the published XSTest labels are the independent check.
- **No certificate of a trained policy in human labels.** Section 10.4 costs one and prepares it; the labels were not collected. Every certificate of a trained policy in this paper is in a guard model’s terms.
- **A rule chosen on the data.** The strata rule of section 8.1 replaced the registered one after the first results on the same pools, and the cells that count as approximate bounds “we use” (section 7.1) were grouped after the results.
- **Scale.** Trained policies are 0.5B to 2B on one consumer card; the frontier evidence is certificates on published traces.
- **One run.** The policy that section 10.4’s samples come from is one training run with one seed, and so is each side of the contrast in section 9.2.
- **Our implementation of others’ methods.** StratPPI’s authors have released no code. Ours matches an independent rewrite of their algorithm and two public libraries’ conventions, which cannot rule out a choice in the authors’ own code that none of these makes. The finding is a finite-sample one and does not contradict the method’s asymptotic guarantee.
- **Validity checks do not prove validity.** A cell within two standard errors of delta is unresolved, and we did not correct for the number of cells.
- **Crossed designs.** The AgentDojo certificate takes the benchmark’s injection tasks as fixed. No bound we specified in advance held when they are treated as sampled.
- **Benchmarks are not deployments.** Sections 10.1–10.2 certify a rate over a benchmark’s task distribution. The clustered bound treats user tasks as sampled from a population of tasks like them; nothing is claimed about tasks unlike them.
- **Approximate bounds.** b1w, the bootstrap-t limits and the cluster bootstrap are checked by resampling, not proved at finite n. Table 4 is the evidence, with its Monte Carlo error.
- **Provenance.** The per-prompt data of the rounds cited from the companion paper did not survive a disk loss; those numbers trace to that paper’s tables and cannot be regenerated without retraining (Appendix D).

12 Related work

Seldonian algorithms. Thomas et al. (2019) define the framework: a candidate/safety split, a high-confidence test, and NSF (sections 2 to 4). We use its safety test unchanged and study what it certifies when the constrained quantity is a judged property of generated text.

Error bars for evaluations. Miller (2024) sets out standard errors for language-model evaluations, including clustered ones. Bowyer et al. (2025) show that normal-approximation intervals fail at small n and recommend alternatives. Our Table 4 agrees and adds the one-sided, selection-robust setting and the failure on crossed designs (section 9.3), which clustered standard errors in one dimension do not cover.

Risk control. Risk-controlling prediction sets (Bates et al., 2021), Learn-then-Test (Angelopoulos et al., 2025) and conformal risk control (Angelopoulos et al., 2024) calibrate a post-processing parameter on held-out data so that a risk is controlled with high probability or in expectation. They share the finite-sample, distribution-free aim and the possibility of finding no valid setting. They control a wrapper around a fixed model; we certify a policy that training has changed, on a label a judge defines. Khosravi and Huo (2026) give anytime-valid selective risk control at deployment for models trained with verifiable rewards; their guarantee is online and per decision, ours is a pre-deployment statement about a rate.

Prediction-powered inference. Angelopoulos et al. (2023a, 2023b) combine a few gold labels with many model predictions; Boyeau et al. (2025) apply it to model evaluation; Fisch et al. (2024) add stratification; Zrnic and Candès (2024) and Gligorić et al. (2025) choose which items to label; Zrnic (2024) gives a bootstrap; Csillag et al. (2025) give e-value versions. Sections 8.1 and 8.2 sit inside this line. What we add is evidence on when the intervals hold at the sample sizes and rates of a safety test (the normal-quantile intervals of PPI++ and StratPPI and PPBoot’s percentile limit are over their level; their estimators with a bootstrap-t limit and proportional allocation are not), a stratifier that a relative constraint supplies for free, a count rule for when to drop the judge, the measured gap between a judge’s verdict and its logit, and the negative of section 9.2: the predictions’ relation to the label is not stable under training that targets the label.

Betting bounds. Waudby-Smith and Ramdas (2024) give the betting confidence sequences we use for bounded means and for the paired bound of section 10.4; the stratified constructions tried in section 8.1 follow Spertus and Stark (2022).

Benchmarks and labels. XSTest (Röttger et al., 2024) and OR-Bench (Cui et al., 2025) supply the benign prompts; XSTest’s three-class annotation scheme is the basis of our refusal guideline. AgentDojo (Debenedetti et al., 2024) supplies the crossed design of section 10.2. Qwen3Guard (Zhao et al., 2025) is the guard.

13 Conclusion

A rate with an error bar is not a certificate. The guarantee belongs to the whole pipeline: how the policy was chosen, how the sample was drawn, and what produced the label. Where we could check against a known truth, no exact bound was over its level, and three common shortcuts were, in their worst cells by wide margins: a normal quantile at small rates, a per-pair bound on a crossed design, and a judge calibration carried to a policy it was not measured on. The last has a practical consequence. Gold labels belong on the responses of the policy being certified, each time the policy changes. What this paper does not contain is a certificate of a trained policy in human labels. We have costed one and fixed its analysis in advance; running it, with more than one annotator, is the next step.

References

Verified against the publisher or arXiv page on 2026-10-04; BibTeX in `reports/paper_certification.bib`.

- Angelopoulos, Bates, Fannjiang, Jordan, Zrnic (2023a). Prediction-powered inference. *Science* 382(6671).
- Angelopoulos, Duchi, Zrnic (2023b). PPI++: Efficient prediction-powered inference. arXiv:2311.01453.
- Angelopoulos, Bates, Fisch, Lei, Schuster (2024). Conformal risk control. ICLR 2024.
- Angelopoulos, Bates, Candès, Jordan, Lei (2025). Learn then test: calibrating predictive algorithms to achieve risk control. *Annals of Applied Statistics* 19(2).
- Bates, Angelopoulos, Lei, Malik, Jordan (2021). Distribution-free, risk-controlling prediction sets. *Journal of the ACM* 68(6).
- Bowyer, Aitchison, Ivanova (2025). Position: Don't use the CLT in LLM evals with fewer than a few hundred datapoints. ICML 2025. arXiv:2503.01747.
- Boyeau, Angelopoulos, Li, Yosef, Malik, Jordan (2025). AutoEval done right: using synthetic data for model evaluation. ICML 2025, PMLR 267:5276–5290. arXiv:2403.07008.
- Csillag, Struchiner, Goedert (2025). Prediction-powered e-values. arXiv:2502.04294.
- Cui, Chiang, Stoica, Hsieh (2025). OR-Bench: an over-refusal benchmark for large language models. ICML 2025 (PMLR 267). arXiv:2405.20947.
- Debenedetti, Zhang, Balunovic, Beurer-Kellner, Fischer, Tramèr (2024). AgentDojo: a dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. NeurIPS 2024 Datasets and Benchmarks.
- Fisch, Maynez, Hofer, Dhingra, Globerson, Cohen (2024). Stratified prediction-powered inference for effective hybrid evaluation of language models. NeurIPS 2024. arXiv:2406.04291.
- Franklin, Schneeweiss, Polinski, Rassen (2014). Plasmode simulation for the evaluation of pharmacoepidemiologic methods in complex healthcare databases. *Computational Statistics & Data Analysis* 72.
- Gligorić, Zrnic, Lee, Candès, Jurafsky (2025). Can unconfident LLM annotations be used for confident conclusions? NAACL 2025. arXiv:2408.15204.
- Khosravi, Huo (2026). Conformal selective acting: anytime-valid risk control for RLVR-trained LLMs. arXiv:2605.20270.
- Miller (2024). Adding error bars to evals: a statistical approach to language model evaluations. arXiv:2411.00640.
- Zhao, Yuan, Huang and 40 others (2025). Qwen3Guard technical report. arXiv:2510.14276.
- Röttger, Kirk, Vidgen, Attanasio, Bianchi, Hovy (2024). XSTest: a test suite for identifying exaggerated safety behaviours in large language models. NAACL 2024. arXiv:2308.01263.
- Spertus, Stark (2022). Sweeter than SUITE: supermartingale stratified union-intersection tests of elections. arXiv:2207.03379.
- Thomas, Castro da Silva, Barto, Giguere, Brun, Brunskill (2019). Preventing undesirable behavior of intelligent machines. *Science* 366(6468).
- Waudby-Smith, Ramdas (2024). Estimating means of bounded random variables by betting. *JRSS-B* 86(1).
- Zrnic (2024). A note on the prediction-powered bootstrap. arXiv:2405.18379.
- Zrnic, Candès (2024). Active statistical inference. ICML 2024 (PMLR 235).

A The StratPPI comparison in detail

The implementation. StratPPI’s authors have released no code. Our implementation agrees with an independent rewrite of their Algorithm 1 to rounding error and reproduces the widths of their own simulation in eight of nine settings (0.231 against 0.222 in the ninth), and the counts of sections 8.1 and 8.2 are the same or worse under the conventions of two public libraries (`ppi_py` and `GLIDE`).

Table A1. Reference-rate strata at delta 0.05: largest miss over checkpoints, and ESS against a random split with a pooled Wilson bound. Five mid-rate labels, two safety-set sizes. The rates in the first column are those of the draws used here; section 8.1 quotes an earlier measurement’s figures for the same labels, whose rates differ by a point and whose ESS is measured from the estimate and not from the truth (section 6.4), so the two sets of ESS differ too (5.13 against 4.75 for refusal of harmful prompts at n_s 100). The ESS is the last checkpoint’s. An ESS beside a bold miss belongs to an arm that is over its level, and the random-split baseline is itself over for the 93% label at n_s 100 (0.057).

label (rate)	n_s	b1w	StratPPI as published	StratPPI estimator, bootstrap-t limit
over-refusal (0.16)	100	0.023; 2.39	0.066 ; 4.11	0.042; 2.58
over-refusal (0.16)	200	0.024; 2.30	0.052; 3.46	0.042; 3.10
over-refusal, pushed by the Lagrangian (0.18)	100	0.024; 2.19	0.056; 3.24	0.036; 2.43
over-refusal, pushed (0.18)	200	0.023; 2.13	0.045; 2.85	0.037; 2.49
non-refusal, encoded prompts (0.09)	100	0.034; 1.38	0.086 ; 2.59	0.045; 1.43
non-refusal, encoded prompts (0.09)	200	0.027; 1.36	0.063 ; 2.11	0.041; 1.59
refusal, harmful prompts (0.65)	100	0.047; 4.75	0.056 ; 6.30	0.033; 1.77
refusal, harmful prompts (0.65)	200	0.040; 5.19	0.045; 6.29	0.034; 5.35
refusal, encoded prompts (0.93)	100	0.054; 0.96	0.032; 0.78	0.031; 0.83
refusal, encoded prompts (0.93)	200	0.045; 1.12	0.026; 0.92	0.033; 1.03

Reference-rate strata (section 8.1). In Table A1, bold marks a miss more than two Monte Carlo standard errors over delta. The predictor is constant in 5 to 7 of the 8 strata. At delta 0.1 StratPPI as published is over its level in 3 of 10 cells. It is also over in every one of the four rare-label cells; **b1w** fails in three of those four. The paper’s own allocation rules do not repair the normal limit: 9 of 26 cells are over with the oracle rule and 11 with the heuristic. With a bootstrap-t limit the estimator is over its level in 6 of 26 cells under the oracle allocation and in 11 under the heuristic, against none of ten under proportional allocation. Under proportional allocation it is ahead of **b1w** in seven cells, by 3–35%, and far behind in one (1.77 against 4.75). A sweep places that drop at a safety set of 100 with 8 or 16 strata; with 4 strata, or with a safety set of 200 at any stratum count, it is absent. A stratified Wald-t limit on the same strata also held in all ten cells (largest miss 0.036, median ESS 2.06). In our predictor comparison the predictor is the same reference rate that defines the strata; splitting the reference samples, four for the strata and four for the predictor, did not change the picture (median ESS 1.97 against 2.10).

Judge-logit strata (section 8.2). StratPPI as published is over its level in 26 of 28 cells under each library’s conventions, in 28 of 28 with the paper’s oracle allocation and in 24 with its heuristic. The cleanest evidence is the judge wording with no constant stratum: 13 of 14 cells over, with misses from 0.056 (the one cell that is not over) to 0.187. With a bootstrap-t limit and the oracle allocation the estimator is over its level in 14 of 28 cells, against none under proportional allocation.

B The mechanism behind section 9.2

Our account of the mechanism, which no experiment here isolates, is a change in the form of the positives. The reward pulled toward long helpful answers and the multiplier penalised the guard’s refusal flag. After training, the responses the guard flags are long refusals with an explanatory, helpful-sounding body: by the guard and a surface pattern (a refusal opener and more than 40 words), their share of the guard’s flags rose from 37% in the reference to 82% in the trained policy. The rubric judge reads that body as an answer: of 93 guard-flagged refusals after training it called 89 answers. Human labels side with the guard. Of 18 such responses of the trained policy, the annotator labelled 12 refusals, 6 answers and none refuse-then-answer, and refuse-then-answer is 3.4% of the reference’s responses against 0.3% of the trained policy’s (section 10.3). We had first read these responses as a refusal sentence followed by the answer; the labels do not support that reading, and the surface pattern is not a label (on XSTest’s completions 56% of responses with it are human full refusals and 34% partial ones).

C The design check of section 10.4

A check of the whole design on the guard’s real scores, with synthetic labels drawn from the guard’s logit at the first sheet’s rates in each band of the logit (a synthetic population with rates of 9.5% and 10.1%, a true difference of +0.006), puts the miss rates of the four limits at 0.001, 0.040, 0.045 and 0.042 for the pool rate against a level of 0.05, and their mean certified margins at 0.065, 0.037, 0.032 and 0.035. When the 490 prompts are redrawn each time, so that the claim is about their source, the four miss rates are 0.001, 0.051, 0.063 and 0.052 (Monte Carlo standard error 0.005); limit (b1) is not built for that case. The exact bound is about twice as wide as the approximate ones: with few discordant pairs in 300, exactness is the larger cost and the guard’s gain is the smaller one. The synthetic labels disagree with the guard independently across the two policies, which weakens both the pairing and the guard, so these margins are on the wide side and those of Table 10 on the narrow side. The guard’s gain in this check is about 1.4 in labels, against 2.5 in Table 10. In the guard’s own terms, on the 3,920 guard-only responses per policy at 256 tokens, the refusal rate is 16.9% for the reference and 17.2% for the trained policy.

D Reproducing the numbers

The numbers in this paper come from the scripts below, apart from those cited from the companion paper and a few recomputed from the data files. The working draft carries a source tag on each number and a table that resolves each tag to a file.

what	script	output
Table 4, every validity cell under one rule	scripts/validity_recount.py	results/paper/validity_recount.md
sections 8.1 and 8.2, the StratPPI comparison	scripts/stratppi_baseline.py	results/paper/stratppi.md
sections 8.1 and 8.2, implementation check, allocations, PPBoot, sweeps	scripts/stratppi_validate.py	results/paper/stratppi_validate.md
section 9.2, the judge against human labels	scripts/judge_on_labels.py	results/labels/refusal/rubric_vs_human.md
section 10.2, AgentDojo	scripts/agentdojo_recheck.py	results/paper/agentdojo_recheck.md

what	script	output
section 10.3, the guard on XSTest	scripts/xstest_guard.py	results/labels/xstest/ analysis.md
section 10.3, the refusal sheet	scripts/refusal_sheet_build.py, scripts/refusal_labels.py	results/labels/refusal/ analysis.md
section 10.4, the label budget	scripts/p9_budget.py	results/paper/p9_budget.md
section 10.4, the prepared certificate and its design check	scripts/p9_sample.py, scripts/p9_ certificate.py	results/labels/p9/design_ check.md
figures	scripts/paper_figures.py	reports/figs/
references checked against arXiv and CrossRef	scripts/bib_check.py	printed report

Results cited from the companion paper on training (the rows of Table 4 on a trained-policy harm pool, and parts of section 7.2) are printed rates whose per-draw data were lost; they cannot be regenerated without retraining.